


PRODUCT & PLATFORM OVERVIEW

The innovation platform for the agentic era.

One instrumented system of record where humans and AI agents plan, build, ship, and learn — governed by roles and approvals, and powered by Evermind: the platform's own self-updating brain, modeled on the hippocampus, neocortex, and limbic system.

- System of Record
- Evermind brain
- Human-in-the-loop
- Self-hosted & open

EVERMIND



- **Neocortex** — routed SSM experts
- **Hippocampus** — write-through memory
- **Limbic system** — affective dynamics

— THE PROBLEM

AI writes the code. Nobody governs the work.



Single-player tools

One dev, one editor. No shared board, no proof of Done.



Ungoverned autonomy

No approval gates, no audit trail, no cost attribution.



Frozen, rented models

Stale on ship. Vendor cloud. Token rent forever.



The agentic era needs a system of record — not another autocomplete.

Four forces converging in 2026.



Agents crossed the capability line

They plan, call tools, and ship multi-step work.



Governance is now mandated

Auditability became a buying requirement.



Token economics demand ownership

Owning your models restores gross margin.



Seats are becoming outcomes

Platforms now manage mixed human+agent teams.

A platform layer across three expanding markets.

TAM

\$300B+

The total spend on AI software, developer tooling, and emerging agentic platforms as work shifts to human+agent teams.

SAM

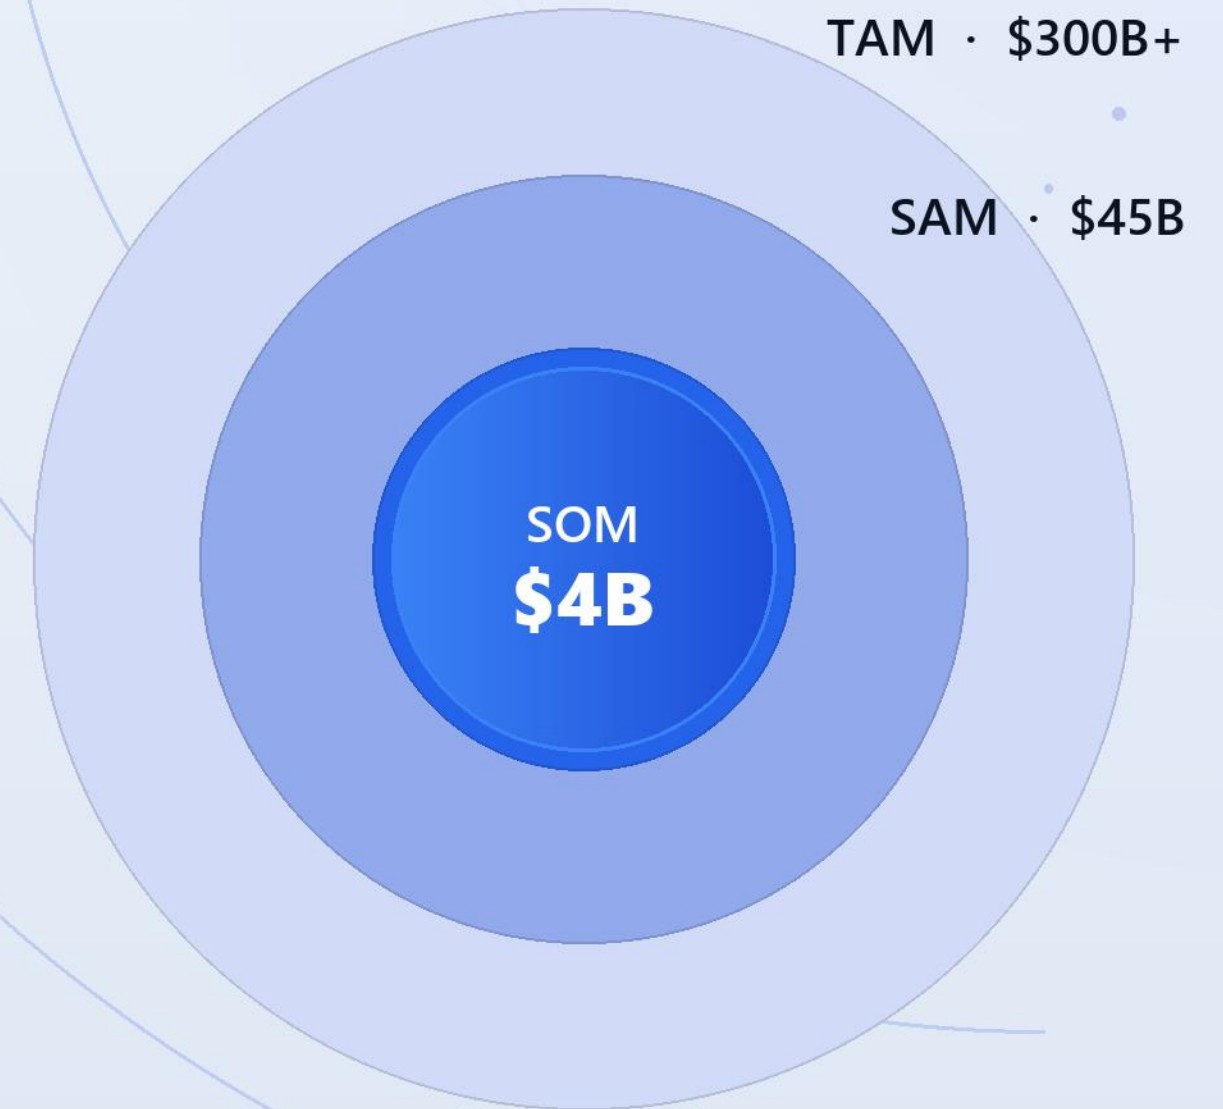
\$45B

Teams buying AI dev platforms, multi-agent orchestration, and the AI-governance layer that makes autonomy auditable.

SOM

\$4B

Our beachhead: engineering orgs that need a self-hosted, governed system of record for a mixed workforce — where we win on ownership.



Figures are illustrative industry estimates (Gartner / IDC / Bloomberg Intelligence ranges), rounded for positioning — not company financials.

— THE SOLUTION

Transformation, delivered on one platform.



One system of record

Every unit of work — human or agent — instrumented and costed.



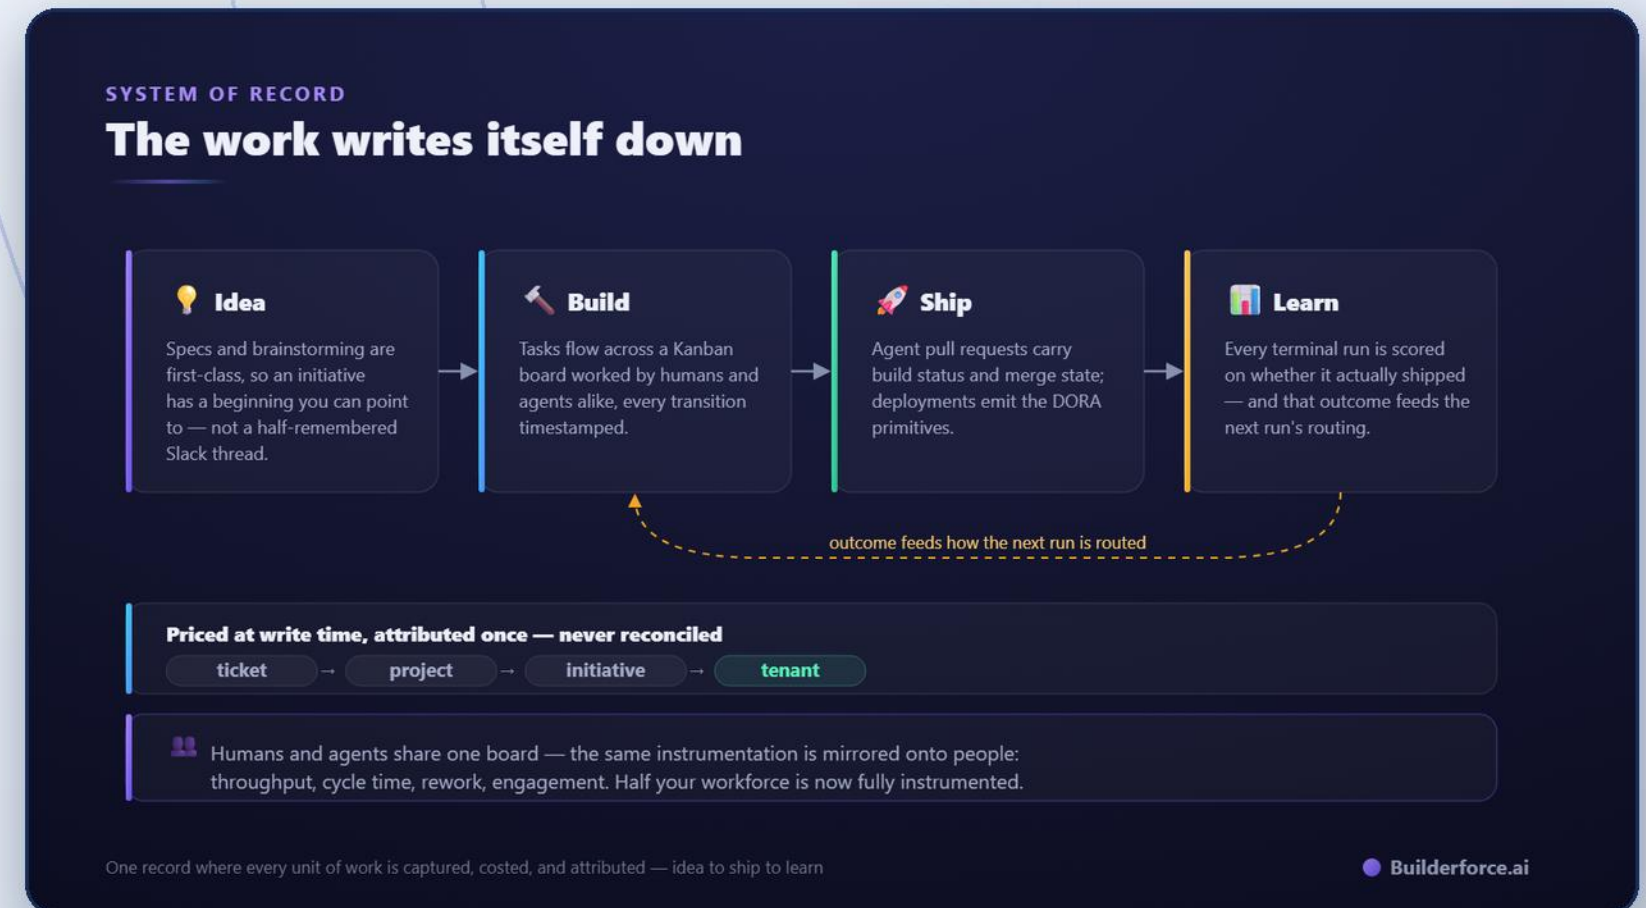
Needs compile into agents

Prose, dataset, chart, or persona becomes a running agent.



A brain that learns

Evermind absorbs every lesson your team ships.



— HOW IT WORKS

Define a need. The system does the rest.



Any input

Language · docs · process chart · persona



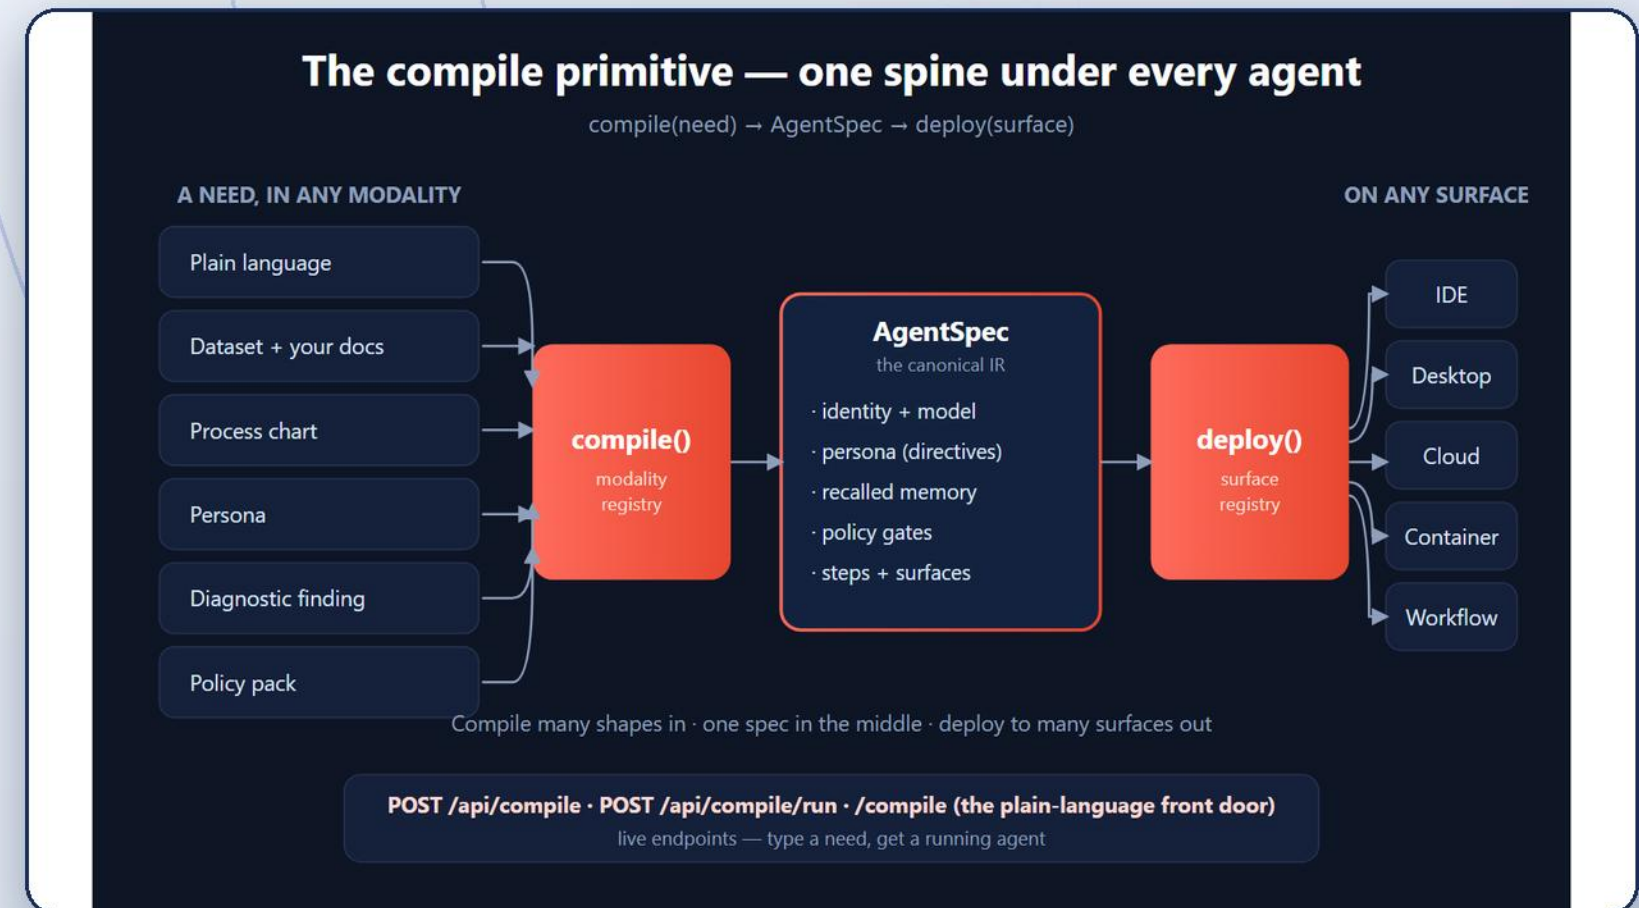
One AgentSpec

Roles, tools, policy, memory — compiled.



Any surface

IDE · desktop · cloud · workflow · MCP



Everything a team needs, in one platform.



Build & Train

- Evermind self-updating model
- Studios: design · voice · LLM · video
- Brain Storm & Dashboard
- In-browser WebGPU training
- Full IDE workspace
- VS Code extension



Orchestrate

- Agentic Workforce Kanban
- Projects, tasks & Gantt
- Planning Spine & Portfolio
- Visual Workflow Builder
- Board connectors (Jira, Linear...)
- Workforce mesh & team chat



Extend

- Knowledge base & SOPs
- Skills & Agent Registry
- Personas & psychometrics
- Prompt library
- Content manager
- MCP — consume & expose



Govern & Operate

- Human-in-the-loop approvals
- Security, RBAC & audit trail
- Observability & cloud delivery
- Quality & error tracking
- Maturity diagnostics
- Tenants & workspaces

AUTONOMOUS WORKFLOWS

A Kanban board that drives itself.



Work lands

A ticket, a crash, or user feedback hits the board.



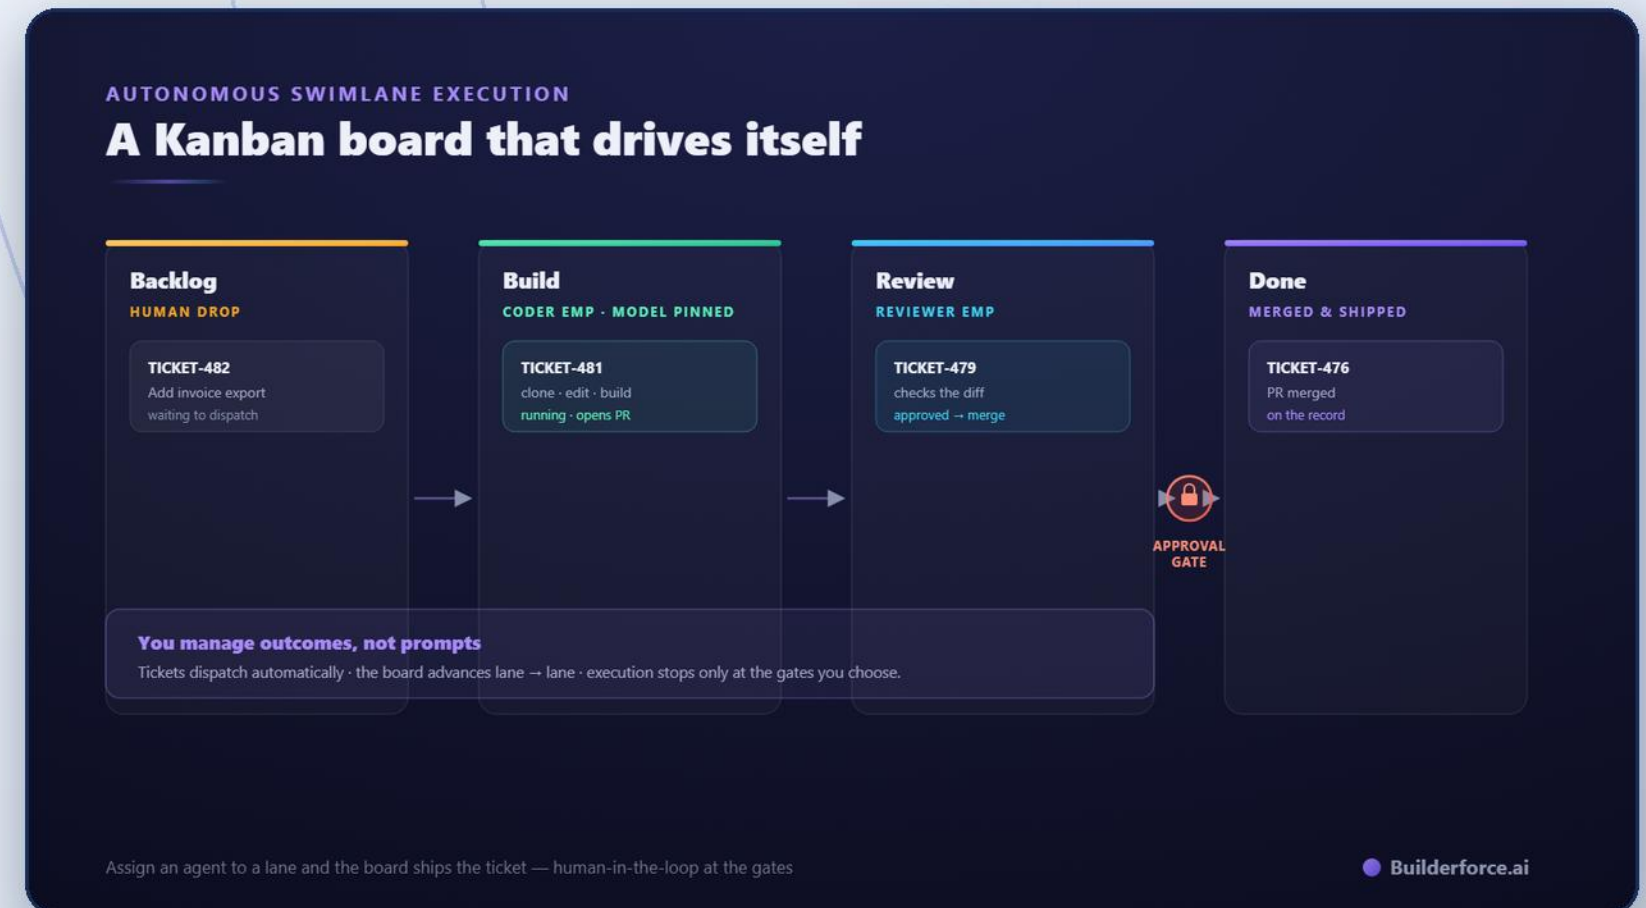
Agents execute

Cloud runners build in real branches and open PRs.



Humans gate

Green CI + your approval before anything merges.



Portfolio to task — one costed spine.



One dated hierarchy

Portfolio → initiative → epic → task, OKRs attached.



Costs roll up

LLM spend + human effort, split CAPEX vs OPEX.



Plans that execute

The same spine agents pick their work from.

THE PLANNING SPINE

One dated hierarchy, cost rolls to every level

Portfolio

\$ rolled total ▲

Initiative

initiative dependencies → critical path

Epic

Task · \$
CAPEX

Task · \$
OPEX

Task · \$
OPEX



Leaf cost rolls up to every ancestor — a real number at any altitude, no parallel finance system, no backfill. One Gantt, all levels dated on the same axis, color-coded by cost class.

CAPEX vs OPEX, resolved

each node carries a cost_class, in priority order:

- 1 · explicit declaration wins
- 2 · inherit from the parent
- 3 · agent classifier by investment category
- 4 · GAAP-conservative default as the floor

⚠ Anomaly flag

child class contradicts parent → sign-off, not silent inherit

📄 CSV export

the whole spine, any date range — same figures, ready for close

Plan, deliver, and account for the same work in the same place — engineering plans, PMO rolls up, finance closes the books

One hierarchy, one set of numbers, no reconciliation step

Builderforce.ai

Frozen models forget. Evermind never does.



Learns as it works

Write-through memory — an update replaces, never drifts.



Distills your frontier models

Frontier quality absorbed into a model you own.



Runs anywhere

Browser, on-device, or inside every agent — \$0 GPU.

WHY IT IS DIFFERENT

A frozen frontier model vs. Evermind

	 FROZEN FRONTIER MODEL	 EVERMIND
Knowledge updates	Frozen at training time — needs a retrain, fine-tune, RAG patch, or hand-edit	Written straight through — an update replaces what came before
Reconciliation	Stale and fresh facts coexist; something must merge them later	Upsert-by-key + invalidate — there is never a reconcile step
Currency	Goes out of date the moment it ships	Never stale — updates land the moment they happen
Footprint	Large; runs only in the vendor cloud	Runs on WebGPU — browser, on-device, or inside every agent
Ownership	Third-party model; a knowledge cutoff you do not control	Yours end to end — open packages, your data

The innovation platform for the agentic era

 Builderforce.ai

WORKS WITH YOUR STACK

Keep the frontier models you already use.



Bring your own keys

30+ providers — Anthropic, OpenAI, Google, local Ollama.



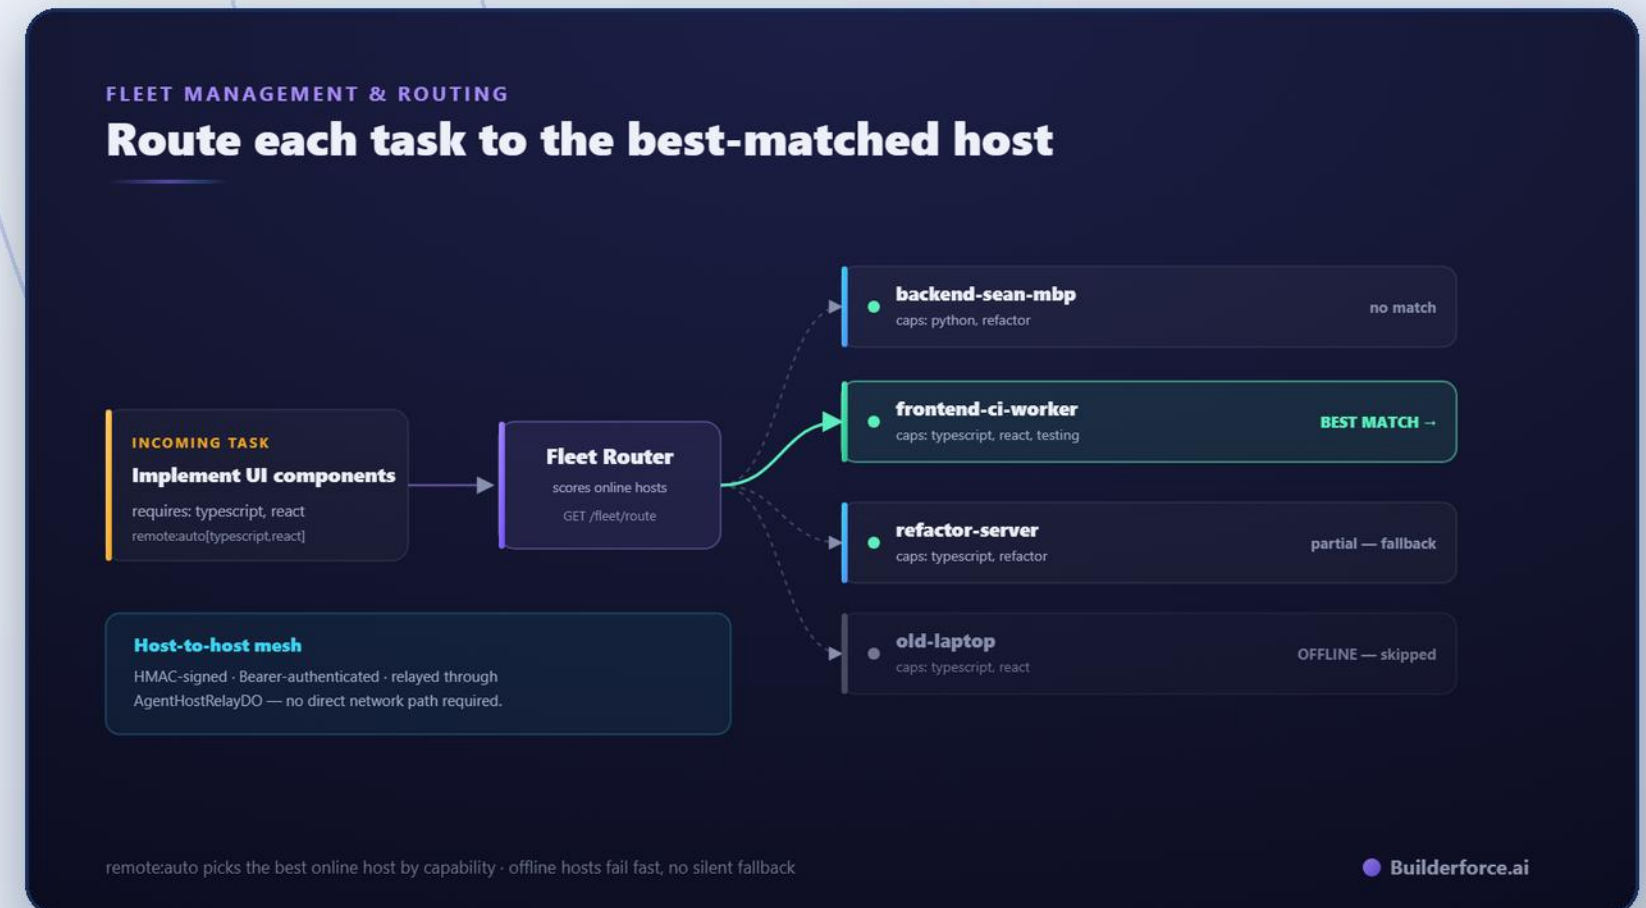
Your Claude subscription

Connect Pro / Max via OAuth. No new contract.



Spend bends down

Learned routing + a semantic cache cut the bill.



Day one: your models and keys. Over time: distilled into a model you own.

Humans and agents. One team, one board.



Role-gated lifecycle

The right reviewer before a ticket advances.



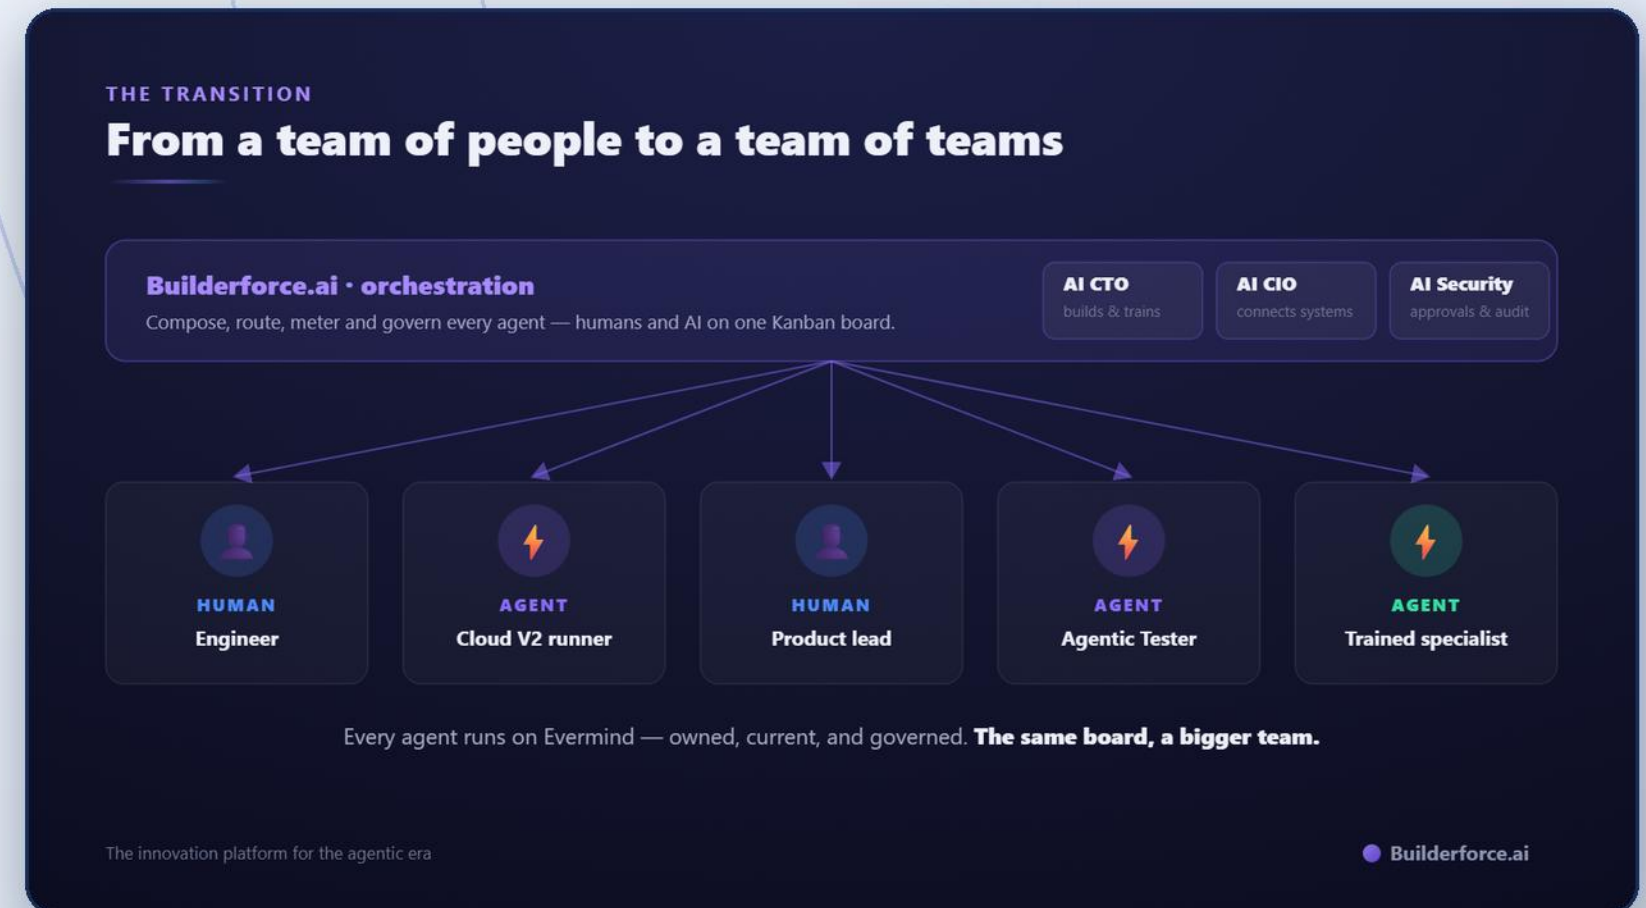
Proof of Done

A Validator agent verifies every item marked Done.



Managers you coach

Typed AI managers, steered by coaching sessions.



Chat, meet, and build — together, live.



Shared threads

Address a person or an @agent in one conversation.



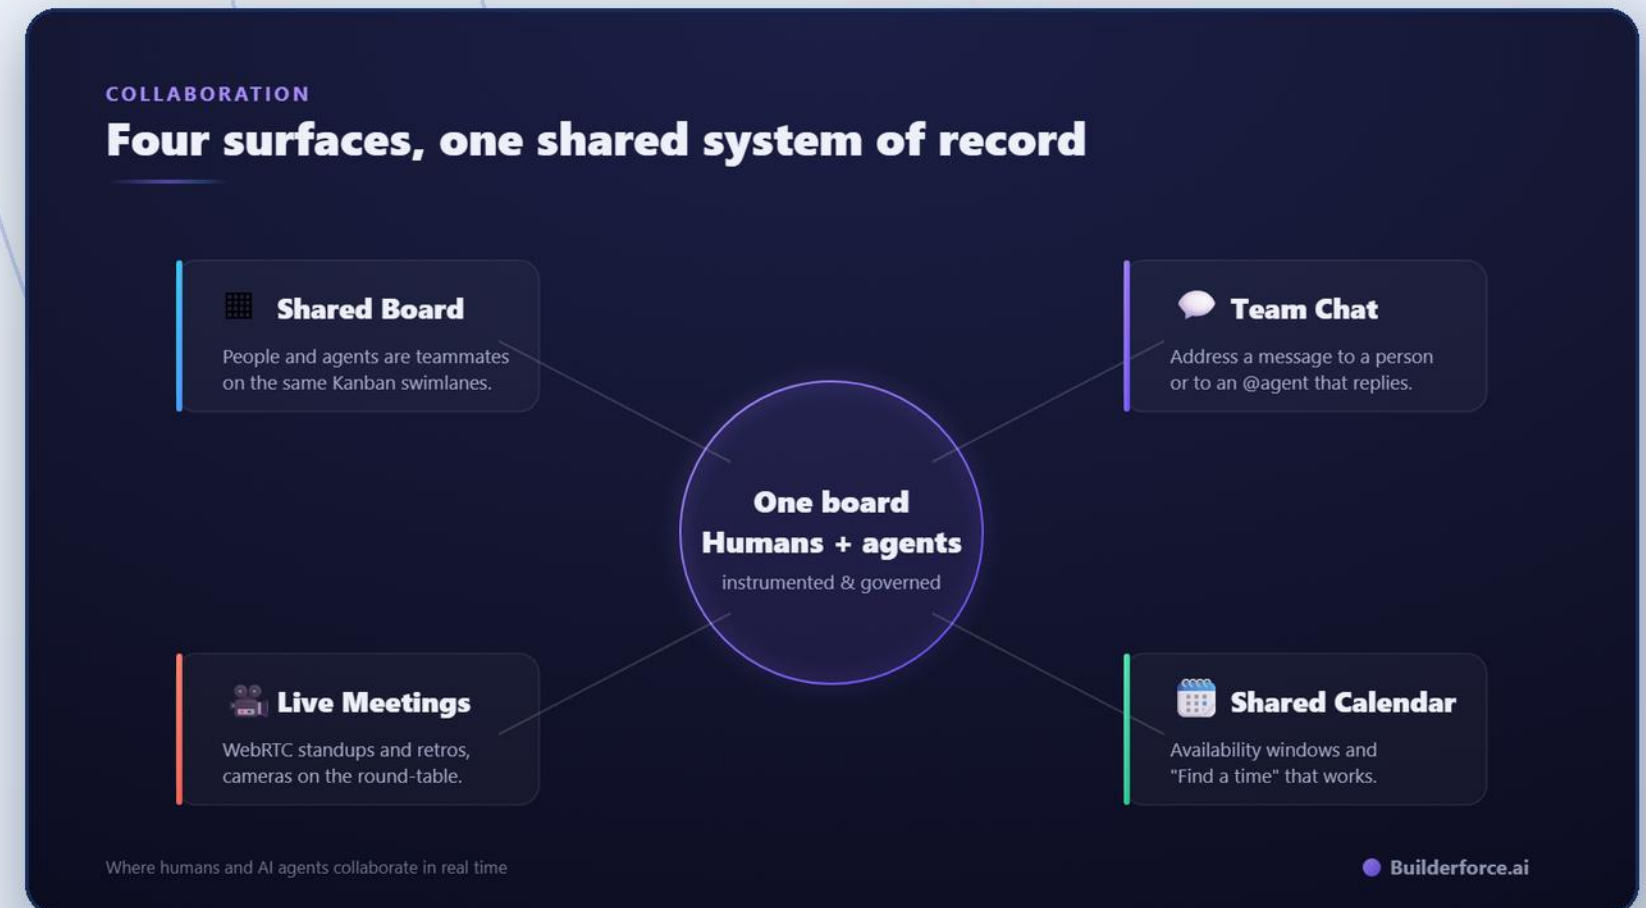
Live meetings

Mesh WebRTC video right on your board.

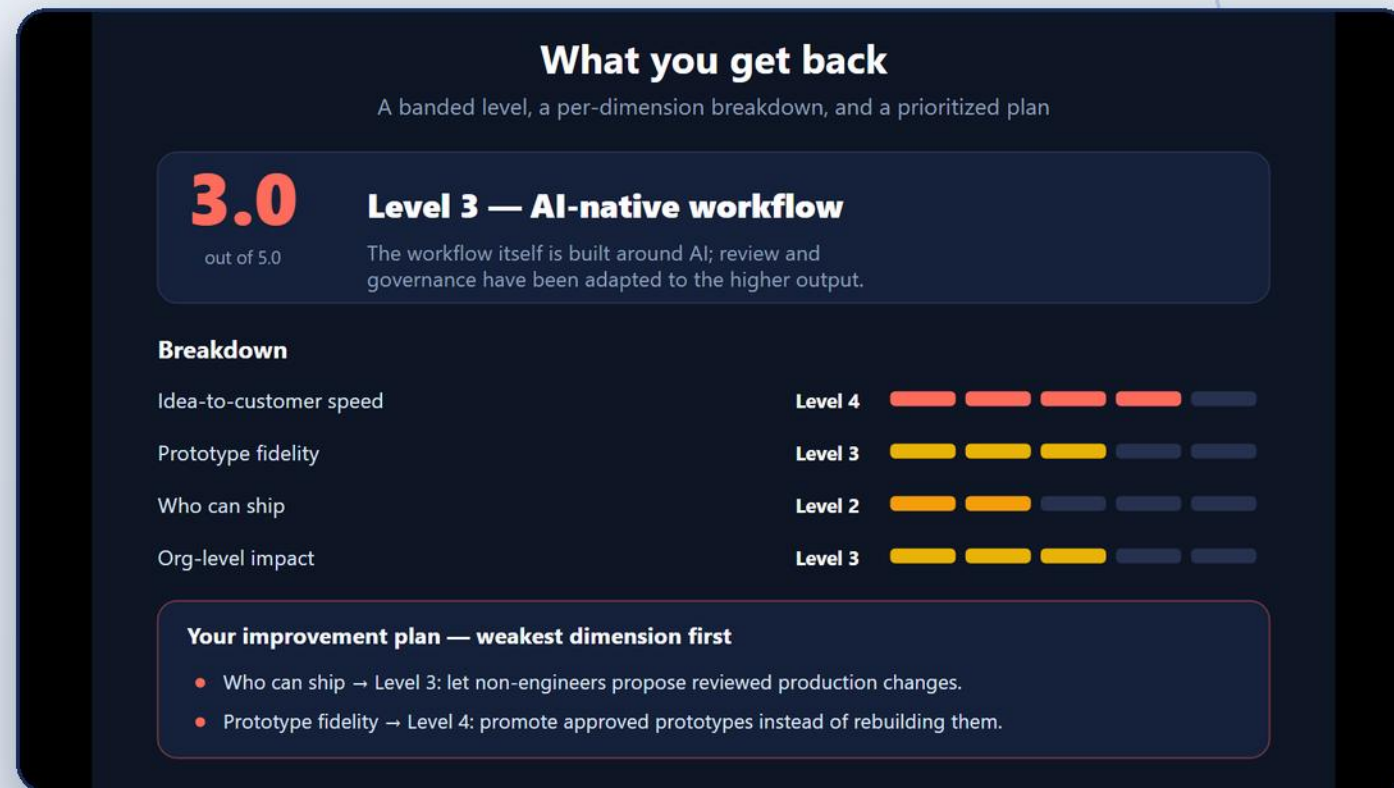


Everywhere you work

Web, VS Code, desktop — the same room.



Every project gets a rating — and a path up.



Maturity diagnostics, project health, and delivery verdicts — measured on shared scales, rolled up to the portfolio.

Autonomy you can put into production.



Approval gates

Destructive actions wait for a human.



Immutable audit

Who, what, at what cost — on the record.



Self-hosted

Air-gapped ready. SOC 2 in progress.



COMPETITIVE LANDSCAPE

Everyone else automates a keystroke. We govern a workforce.



The capabilities single-agent tools structurally lack.

Capability	Builderforce	Copilot	Cursor	Claude Code	Devin	OpenHands
Self-hosted / open source	✓	×	×	×	×	✓
Air-gapped / private deploy	✓	×	×	×	×	✓
Trains your own model on your work	✓	×	×	×	×	×
Self-updating model per project	✓	×	×	×	×	×
Learned, outcome-based routing	✓	×	×	×	×	×
Validator agent — proof of Done	✓	×	×	×	×	×
Role-gated board + audit trail	✓	×	×	×	—	×
Multi-agent orchestration	✓	×	×	×	×	×
Human-in-the-loop approval gates	✓	×	×	×	—	×
Any model provider (30+)	✓	×	—	×	×	✓

Land free. Expand on outcomes, not seats.

Free

\$0 /forever

- ✓ Self-updating model on-device
- ✓ WebGPU training + dataset gen
- ✓ 5 projects · 10K tokens/day
- ✓ 1 agent host

Pro

Most popular

\$29 /seat / mo

- ✓ Everything in Free
- ✓ Unlimited projects · 1M tok/day
- ✓ Approval workflows + audit
- ✓ Fleet mesh + custom roles

Teams

\$20 /seat / mo

- ✓ Everything in Pro
- ✓ Unlimited hosts · 5M tok/day
- ✓ Shared approval inbox
- ✓ Per-seat cost controls

Expansion revenue beyond seats

Managed hosting
\$49 / agent / mo

Agent & template marketplace
Listing + transaction fee

Enterprise license
On-prem Docker, annual

B2B AI gateway
Metered LLM margin

A platform already built — not a pitch of one.

40+

Shipped platform capabilities

350+

Database migrations in production

30+

Model providers supported

3

Run surfaces: IDE · desktop · cloud

2B

On-device params, WebGPU

4

Generative studios: design·voice·LLM·video

Runs where you do

VS Code · desktop · cloud

Own your model

on-device WebGPU, \$0 GPU

Ships real work

crash & ticket → pull request

Deploy anywhere

self-hosted · air-gapped ready

Every figure above is a verifiable product fact — the platform is built, shipping, and self-hostable today.

Open-source led. Bottom-up. Enterprise-ready.



01

Adopt

Free, self-hostable, MIT-spirited. Land via the VS Code Marketplace, GitHub, and the open MCP ecosystem — a single developer starts in minutes.



02

Activate

Teams put humans + agents on one board, wire GitHub / Jira / Linear, and feel the governance and cost visibility no single-player tool gives them.



03

Expand

Seats convert to Pro / Teams; managed hosting, the agent & template marketplace, and the metered AI gateway grow spend with usage.



04

Enterprise

On-prem license, SOC 2, air-gapped deploy, and role-gated governance close regulated buyers where incumbents structurally can't follow.



Distribution already live: the AI gateway powers partner brands (hired.video, burnrateos.com) — the platform sells itself into adjacent products.

A credible path from platform to \$1M ARR.

Q3 2026

\$50K ARR

Teams plan + AI gateway; managed hosting beta; first-agent-in-5-min funnel

Q1 2027

\$500K ARR

Enterprise expansion + marketplace GA; governance-auditor + policy packs

Q4 2026

\$200K ARR

Enterprise pilots + marketplace monetization; SOC 2 Type I; on-prem license

Q2 2027

\$1M ARR

Scale + multi-year enterprise contracts; agent mesh at scale

ARR figures are company targets from the internal roadmap. Product milestones are scoped and in progress today.

Built by a 25-year technology operator.



Sean Hogg

Founder · CTO & Technology Leader

25 years driving business value through technology — CTO, CIO and VP of Product & Technology across mortgage, automotive and SaaS, plus a founder exit. A hands-on builder who ships revenue and cuts cost, now channelling advanced AI & ML into one platform. Dual degrees in Computer Science and Psychology — the foundation for Builderforce's persona + limbic model.

CTO · CIO · VP

25 yrs · AI & ML

CS + Psychology

Founder (exit)

Selected track record

\$6M+

New-product revenue delivered on a single framework in 2 years.

\$1.79M

Operating cost removed through IT modernization & renegotiation.

200+

Features shipped in 8 months leading onshore/offshore teams.

+175%

Project completion lift while cutting spend from \$4M to \$1.7M.

90%

Attrition reduction by building a performance-first culture.



GET STARTED

Build your agentic workforce **on a platform you own.**

Plan, build, ship and govern with a workforce of humans and AI agents on one instrumented board — free to start, self-hostable, and powered by a model that learns as it works.

Start free

\$0 forever · no credit card

Self-host it

open · air-gapped ready

Book a demo

see it on your workflow

builderforce.ai · seanhogg@builderforce.ai