

The innovation platform for the agentic era.

One instrumented system of record where a workforce of humans and AI agents plans, builds, ships, and learns — governed by roles, approvals, and a full audit trail, powered by a model you own that never goes stale.

- System of Record
- Self-updating model
- Human-in-the-loop
- Self-hosted & open

— THE PROBLEM

AI writes the code. Nobody governs the work.

Today's AI dev tools make one developer faster in one editor. None of them make an enterprise's agentic work governable, costed, or trustworthy in production.



Single-player tools

Copilot, Cursor, and Claude Code accelerate one dev in one window. No shared board, no roles, no proof that "Done" is actually done, no memory of the team's work.



Ungoverned autonomy

Agents act with no approval gates, no audit trail, and no cost attribution. Enterprises can't put autonomous work into production they can't review, price, or prove.



Frozen, rented models

Frontier models go stale the moment they ship, send your data to a vendor cloud, and bill by the token forever — with every retrain a bolt-on you don't own.



The agentic era needs a system of record — not another autocomplete.

Four forces converging in 2026.



Agents crossed the capability line

Frontier models can now plan, call tools, and ship multi-step work. The bottleneck moved from capability to coordination, trust, and cost.



Enterprises are mandating AI — and governance

Boards want AI adoption and auditability at once. “Who did what, at what cost, approved by whom” is now a buying requirement, not a nice-to-have.



Token economics demand ownership

Usage-based bills scale with success. Routing, caching, and owning a model that runs on-device are how AI-native orgs keep gross margin.



Dev tools shift from seats to outcomes

The market is moving from per-seat autocomplete to outcome-priced platforms that manage a mixed workforce of people and agents.

MARKET OPPORTUNITY

A platform layer across three expanding markets.

TAM

\$300B+

The total spend on AI software, developer tooling, and emerging agentic platforms as work shifts to human+agent teams.

SAM


\$45B

Teams buying AI dev platforms, multi-agent orchestration, and the AI-governance layer that makes autonomy auditable.

SOM

\$4B

Our beachhead: engineering orgs that need a self-hosted, governed system of record for a mixed workforce — where we win on ownership.



Figures are illustrative industry estimates (Gartner / IDC / Bloomberg Intelligence ranges), rounded for positioning — not company financials.

— THE SOLUTION

Builderforce.ai is the **system of record for the agentic enterprise.**

Every unit of work — human or agent — is instrumented, costed, and attributed from idea to ship to learn, on one board governed by roles, approvals, and a full audit trail.



01

Instrument every unit of work

Plan, assign, and track people and agents on one live board — portfolio to task, with cost and DORA rolling up automatically.



02

Compile a need into an agent

Define a need in any modality — prose, a dataset, a process chart, a persona — and the platform compiles it into an agent that runs anywhere.



03


Own a model that learns

Evermind is your self-updating model: it learns from your team's work, distills frontier quality into a model you run on-device.

— HOW IT WORKS

Define a need. The system compiles and runs the agent.

`compile( need, modality )` → AgentSpec → `deploy( AgentSpec, surface )`



One spine, many inputs, many surfaces — the same compiled agent runs in the IDE, the desktop, the cloud, or a workflow step.

One system. Six capability pillars.



Plan & Deliver

- Planning spine: portfolio → task
- Kanban, OKRs, Gantt & calendar
- Cost & DORA roll-up, no reconcile



Evermind Model

- Self-updating, on-device model
- Distill any frontier model in
- Learned routing + generative media



Agentic Workforce

- Humans + agents on one board
- Role-gated, sign-off proof of Done
- AI managers you coach



Govern & Trust

- Approval gates + immutable audit
- RBAC, per-tenant, AES-256
- Self-hosted / air-gapped ready



Operate

- Incidents, on-call & monitoring
- Error → fix → pull request
- Native GitHub App + CI gates



Collaborate

- Multi-party human+agent chat
- Live WebRTC meetings & calendars
- Knowledge base, all in VS Code

A model that learns as it works — and never goes stale.

Frozen frontier models go out of date the moment they ship, and every update is a bolt-on retrain you don't own. Evermind writes new knowledge straight through, so an update simply replaces what was there.

Currency

Never stale

Knowledge lands in the model the moment it happens — no retrain cycle, no reconciliation step.

Footprint

Runs anywhere

In the browser, on-device, or embedded in every agent via WebGPU — zero cloud GPU bills.

Ownership

Yours end-to-end

Open packages, your data, your weights — no third-party model dependency or per-token rent.



Write-Through Cognition

the law that governs the brain

- **Distill a frontier teacher**

Pin Opus, Mistral or GLM as a teacher; your model absorbs frontier quality from real work you ship.

- **Shared-expert hybrid**

A dense online backbone + lazily-loaded routed SSM experts — specialist depth without a frozen blob.

- **Write-through memory**

Every fact upserts by a stable key and invalidates its old recall — the loop corrects in place, never drifts.

- **Generative media**

The same engine now generates images and video on-device, not just text.

Humans and AI agents on one board.

Staff each project from a first-class role taxonomy — people and agents in the same roster. Quality is gated at the board, not hoped for after the fact.



Role-gated lifecycle

Swimlanes require the right reviewer before a ticket advances. Every ticket carries a participation manifest — the roles it needs, resolved to the humans and agents capable of each.



Validator — proof of Done

A built-in agent reviews every item marked Done, records a verdict, and opens GAP tasks for anything short. Acceptance reviews land right on the pull request.



AI managers you coach

Typed managers — Dev, QA, Service-Desk, DevOps — run their reports. When one gets it wrong, a human runs a Coaching Session that steers it going forward.



Hire people and agents

Bring on vetted freelance talent (résumés via hired.video) billed by tracked time, alongside a registry of trained agents you hire and call from inside your own agent.

Autonomy you can put into production.

The controls that turn “agents that write code” into “agents an enterprise can approve, price, and prove” — built into the platform, not bolted on.



Human-in-the-loop gates

Every dispatch passes one shared approval gate; destructive and cross-boundary actions suspend for a person to approve.



Immutable audit trail

One unified activity log records who, when, verdict, and contribution per role — append-only, evidence-ready.



RBAC & isolation

Role-based access, default-deny sign-off, per-tenant isolation, and AES-256-GCM encrypted credentials.



Self-hosted & air-gapped

Run the whole workforce on-prem with local models — no code or prompt leaves your network. SOC 2 Type I in progress.



Cost attribution & FinOps

LLM spend priced at write time + human effort roll up per task, split CAPEX vs OPEX with anomaly flags.



Proof of Done

Role-gated sign-off + Validator verdicts mean “Done” carries evidence, not just a status change.

COMPETITIVE LANDSCAPE

Everyone else automates a keystroke. We govern a workforce.



— WHY WE WIN

The capabilities single-agent tools structurally lack.

Capability	Builderforce	Copilot	Cursor	Claude Code	Devin	OpenHands
Self-hosted / open source	✓	×	×	×	×	✓
Air-gapped / private deploy	✓	×	×	×	×	✓
Trains your own model on your work	✓	×	×	×	×	×
Self-updating model per project	✓	×	×	×	×	×
Learned, outcome-based routing	✓	×	×	×	×	×
Validator agent — proof of Done	✓	×	×	×	×	×
Role-gated board + audit trail	✓	×	×	×	—	×
Multi-agent orchestration	✓	×	×	×	×	×
Human-in-the-loop approval gates	✓	×	×	×	—	×
Any model provider (30+)	✓	×	—	×	×	✓

Land free. Expand on outcomes, not seats.

Free

\$0 /forever

- ✓ Self-updating model on-device
- ✓ WebGPU training + dataset gen
- ✓ 5 projects · 10K tokens/day
- ✓ 1 agent host

Pro

Most popular

\$29 /seat / mo

- ✓ Everything in Free
- ✓ Unlimited projects · 1M tok/day
- ✓ Approval workflows + audit
- ✓ Fleet mesh + custom roles

Teams

\$20 /seat / mo

- ✓ Everything in Pro
- ✓ Unlimited hosts · 5M tok/day
- ✓ Shared approval inbox
- ✓ Per-seat cost controls

Expansion revenue beyond seats

Managed hosting
\$49 / agent / mo

Agent & template marketplace
Listing + transaction fee

Enterprise license
On-prem Docker, annual

B2B AI gateway
Metered LLM margin

A platform already built — not a pitch of one.

40+

Shipped platform capabilities

350+

Database migrations in production

30+

Model providers supported

3

Run surfaces: IDE · desktop · cloud

2B

On-device params, WebGPU

4

Generative studios: design·voice·LLM·video

EARLY GO-TO-MARKET SIGNAL

Design partners

Waitlist / signups

\$ ARR run-rate

VS Code installs

Metrics in the top grid are verifiable product facts; the bottom strip is founder-completed go-to-market data.

Open-source led. Bottom-up. Enterprise-ready.



01

Adopt

Free, self-hostable, MIT-spirited. Land via the VS Code Marketplace, GitHub, and the open MCP ecosystem — a single developer starts in minutes.



02

Activate

Teams put humans + agents on one board, wire GitHub / Jira / Linear, and feel the governance and cost visibility no single-player tool gives them.



03

Expand

Seats convert to Pro / Teams; managed hosting, the agent & template marketplace, and the metered AI gateway grow spend with usage.



04

Enterprise

On-prem license, SOC 2, air-gapped deploy, and role-gated governance close regulated buyers where incumbents structurally can't follow.

Distribution already live: the AI gateway powers partner brands ([hired.video](#), [burnrateos.com](#)) — the platform sells itself into adjacent products.

A credible path from platform to \$1M ARR.



ARR figures are company targets from the internal roadmap. Product milestones are scoped and in progress today.

Built by an operator — and by the product itself.



Sean Hogg

Founder & Chief Engineer

Full-stack founder-engineer who designed and shipped the entire platform — the self-updating model, the agentic workforce, the governance layer, and the multi-surface runtime — as a human-in-the-loop agentic workforce. The product is the proof: Builderforce.ai builds Builderforce.ai.

Platform architecture

Applied ML / SSM

Enterprise SaaS

Developer tooling

What we're building next

- **Founding GTM / DevRel**
Turn open-source adoption into a repeatable enterprise motion.
- **Applied ML engineer**
Push Evermind's distillation + on-device generative quality.
- **Enterprise CSM / solutions**
Land and expand regulated design partners.
- **Design partners (now)**
Co-build the governed agentic workflow for your org.



THE ASK

Own the system of record for the agentic enterprise.

We're raising a seed round and onboarding design partners to bring a governed workforce of humans and AI agents — powered by a model you own — into production teams.

Seed capital

\$__ to scale GTM + ML

Design partners

Regulated eng orgs, now

Intros

Enterprise AI buyers

builderforce.ai · seanhogg@builderforce.ai